Improved eukaryotic detection compatible with large-scale automated analysis of metagenomes

Wojtek Bazant1*, Ann Blevins2, Kathryn Crouch1#, Daniel P. Beiting2*#

1Institute of Infection, Immunity and Inflammation, College of Medical, Veterinary and Life Sciences, University of Glasgow, United Kingdom

2Department of Pathobiology, School of Veterinary Medicine, University of Pennsylvania, Philadelphia, Pennsylvania, 19104, USA

*To whom correspondence should be addressed. E-mail: beiting@upenn.edu and
#Indicates co-senior authors

Keywords: metagenome, shotgun metagenomics, eukaryotes, bioinformatics, fungi

Abstract

Background

Eukaryotes such as fungi and protists frequently accompany bacteria and archaea in microbial communities. Unfortunately, their presence is difficult to study with shotgun sequencing techniques since prokaryotic signal dominates in most environments. Recent methods have made detection of known eukaryotes possible using eukaryote-specific marker genes, yet they do not incorporate strategies to detect the presence of unknown eukaryotes.

Results

Here we present CORALE (for Clustering of Reference Alignments), a method for identification of eukaryotes in shotgun metagenomic studies that uses the previously reported EukDetect marker gene reference but increases sensitivity of detection through multiple alignments and Markov clustering. Using a combination of simulated datasets and large publicly available human microbiome studies, we demonstrate that our method not only improves sensitivity, but is also capable of inferring the presence of eukaryotes not included in the marker gene reference, such as novel species and strains. We then deploy CORALE on our MicrobiomeDB.org resource, demonstrating adequate reliability and throughput.

Conclusion

CORALE allows eukaryotic detection to be automated and carried out at scale. Since our approach is independent of the reference used, it is applicable to other contexts where shotgun metagenomic reads are matched against redundant but non-exhaustive databases, like taxonomic classification of viral reads.

Background

Eukaryotic microbes are a large and phylogenetically diverse group of organisms that includes both pathogens and commensals, the latter of which are emerging as important modulators of health and disease. Protists include many important pathogens of humans and other animals, such as Cryptosporidium, Toxoplasma, Eimeria, Trypanosoma, and Plasmodium. Many fungi are also well-studied pathogens affecting a diverse range of hosts. For example, *Aspergillus fumigatus* is an important cause of respiratory disease in humans ("Aspergillus Fumigatus and Aspergillosis" 1999); *Magnaporthe oryzae* is the most important fungal disease of rice globally (Wilson and Talbot 2009); while *Pseudogymnoascus destructans* is the cause of White-Nose Syndrome, one of the most devastating diseases of bats (Wibbelt et al. 2010). However, recent data also suggests that non-pathogenic commensal fungi are significant as modulators of the human antibody repertoire Ost et al. (2021); intestinal barrier integrity, and colonization resistance Jiang et al. (2017). This diverse array of host-microbe interactions and host phenotypes influenced by eukaryotic microbes underscores the importance of studying this class of organisms in their natural habitat. Unfortunately, the ability to carry out culture-independent analysis of eukaryotic microbes is severely hindered by their low abundance relative to bacteria, which makes accurate detection a challenge and means that these organisms are commonly overlooked in metagenomic studies (Laforest-Lapointe and Arrieta 2018). For example, an analysis of stool metagenomes in healthy adults participating in the Human Microbiome Project (Consortium et al. 2012) showed only 0.01% reads aligning to fungal genomes(Nash et al. 2017).

Several methods have been developed to improve the detection of eukaryotes in complex samples. Targeted sequencing of internal transcribed spacer regions (ITS) (Schoch et al. 2012) is a common approach but prevents simultaneous profiling of other members of the microbiome. Alternatively, collections of curated fungal genomes have been successfully used for strain-level identification of Blastocystis from stool (Beghini et al. 2017), however, pitfalls associated with non-specific or erroneous parts of reference genomes (R Marcelino, Holmes, and Sorrell 2020) together with computational challenges associated with carrying out alignments to very large collections of reference genomes Breitwieser, Baker, and Salzberg (2018), mean that alternative approaches are needed to enhance the discovery of eukaryotes from the vast amount of metagenomic data already available in the public domain. One attractive solution to this challenge was recently proposed in important work by Lind and Pollard (Lind and Pollard 2021), who base their method for sensitive and specific identification of eukaryotes in metagenomic studies, EukDetect, on alignments to over 500,000 universal, single-copy eukaryotic marker genes.

We recently sought to add EukDetect results our web-based resource, MicrobiomeDB.org (Oliveira et al. 2018), in order to allow eukaryote detection across a range of human metagenomic studies currently available on the site. Since the

EukDetect pipeline does not allow for adjustment of filtering thresholds and it is not packaged for containerised deployments, we decided to implement our own tool, with a more flexible software architecture. We decided to retain EukDetect's reference of marker genes with the aim of producing directly comparable results, kept `bowtie2` (Langmead and Salzberg 2012) since it has been shown to be a sensitive aligner (Thankaswamy-Kosalai, Sen, and Nookaew 2017), and evaluated different filters used by EukDetect. We noticed that a filter based on MAPQ scores, whose function is to remove uncertain hits, also removes good alignments. Studying MAPQ scores in simulated data has shown us that alignments to marker genes reflect relations of each source of reads to parts of the reference that are most similar to it. This led us to develop CORALE (for Clustering of Reference Alignments), an approach to processing marker gene alignments based on exploiting information in shared alignments to reference genes through clustering. CORALE is able to sensitively detect eukaryotes from the reference while also enabling inference of novel species not present in the reference.

Results

Species-specific impact of MAPQ filtering

We base our study of MAPQ scores in simulated data on the `wgsim` (Li 2011) software for sampling reads, `bowtie2` used like in EukDetect for aligning simulated reads, and a custom Python program for recording properties of alignments like MAPQ together with degree of mismatch between source taxon of each read and the taxon of the sequence it aligns to. To quantify outcomes, we calculate the proportion of sampled reads that correctly map (recall), as well as the correctly mapped reads as a proportion of total reads that map to any reference (precision) (Meyer et al. 2019).

Not surprisingly, when reads in a metagenomic sample exactly match the reference (as simulated by sampling from the reference itself) they are accurately mapped to the correct taxon with a precision and recall of 95.1%. Applying a MAPQ >= 30 filter increases precision to 99.7% and decreases recall to 91.7%. This translates to 8% of reads mapping with MAPQ < 30 with just under half of those being incorrectly mapped. In contrast, with MAPQ >= 30 only 0.3% are incorrectly mapped.

Stratifying these values by the source taxon of the reads reveals a structural component to the difficulty of mapping the reads, as well as the efficacy of the MAPQ >= 30 filter (Figure 1). For example, out of 3977 taxa whose reads map back to the reference, reads from 1908 map with 100% precision (Figure 1, red points). After applying the MAPQ >= 30 filter, 1105 more taxa map with 100% precision. Despite this clear improvement after applying the MAPQ filter, 146 taxa still map with precision lower than the pre-filter overall total of 95.1% (Figure 1, dotted line). This includes numerous species of Aspergillus (Figure 1A), Leishmania (Figure 1B), and Trichinella (Figure 1C), all of which are important pathogens of humans and other mammals. Furthermore, applying the MAPQ filter results in decreased precision for five taxa, including the fungi *Fusarium cf. fujikuroi NRRL 66890* and *Escovopsis sp. Ae733*, as well as the protists *Favella ehrenbergii*, *Leishmania peruviana*, and *Mesodinium rubrum*.

Since the diversity of eukaryotic taxa far exceeds the currently discovered species, let alone species present in the EukDetect reference (Mora et al. 2011), we then modify this experiment to study the possibility of detecting 'novel' species. We split the species level markers in the EukDetect reference into a hold-out set of 371 taxa, and a reference set of 3343 remaining taxa. Sampling from the hold-out set and aligning to an index built from the reference set is more successful without applying the MAPQ >= 30 filter: same-genus precision and recall 82% and 30% without the filter, comparing to 83.6% precision and a much-diminished recall of 7% after applying the filter. Source taxon is a structural component here, too: while applying the MAPQ >= 30 filter increases the number of taxa which only map to the correct genus from 48 to 152, it also increases the number of taxa which don't map at all from 49 to 175.

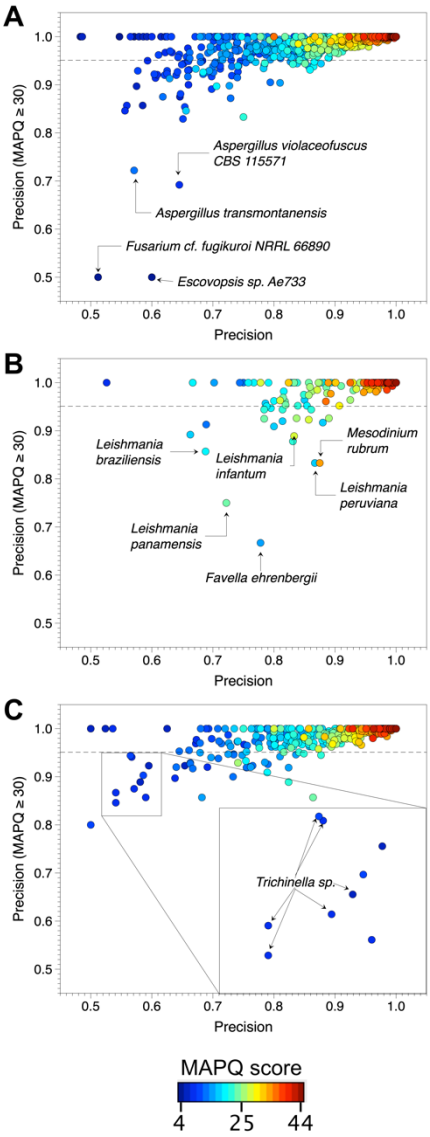

Figure 1: **The improvement from the MAPQ >= 30 filter varies by species.**

We are also able to approximate a possiblity of 'novel' and non-reference species by sampling reads from the reference, and then mutating them using the `wgsim` mutation rate model. Starting from non-modified reads, and increasing the mutation rate in increments, makes the alignment problem increasingly harder. Over the [0, 0.200) range of `wgsim` mutation rate parameter M, recall declines from 95.1% to below 10% while precision stays between 95-96% for all reads and >= 99% for reads with MAPQ >= 30, consistently with bowtie2 preserving precision over recall [16]. Applying the MAPQ >= 30 filter causes recall to drop much faster: for example, when M = 0.1, recall is 68.3% without the MAPQ >= 30 filter and 5.0% with the filter.

CORALE leverages Markov clustering for reference-based eukaryote detection

We wrote our software, CORALE - for Clustering of Reference Alignments - as a Nextflow workflow wrapping a Python module. It provisions sequence files, aligns them to the refernce of markers, and produces a taxonomic profile through a seven-step procedure (Figure 2). First, we run `bowtie2` and keep all alignments that are at least 60 nucleotides in length (Figure 2, step 1), to ensure that the matches contain enough information to be marker-specific, and then use Markov Clustering (MCL) with counts of shared alignments between marker genes to produce marker clusters (Figure 2, step 2). We then compute average % match identity for each marker gene, as well as an average match identity for all alignments within each marker cluster (Figure 2, step 3), for the purpose of rejecting any taxon with <= 50% of marker genes having % identity that is lower than the average % identity for the cluster (Figure 2, step 4). We then group remaining taxa into taxon clusters using MCL with counts of multiply aligned reads (Figure 2, step 5), which allows us to reflect ambiguity of identification in reporting the hits. We report unambiguous matches (defined as having average alignment identity of at least 97%, two different reads aligned to at least two markers) as is (Figure 2, step 6), remove other taxa in clusters where there were any unambiguous and then consider the remaining taxon clusters to each be one potential hit, reporting it if it's a strong ambiguous match (defined as having at least four markers and eight reads) by joining names of taxa in the cluster and prepending with a "?" (Figure 2, step 7).

In the CORALE Python module, all thresholds are adjustable, but we recommend above values based on our evaluation on simulated and human microbiome data. The software also supports a few other filters developed in the course of our experimentation, like a filter based on the fraction of primary to secondary alignments for each taxon which is a simpler but slightly inferior alternative to the filter based on marker cluster averages. CORALE has rich reporting capabilities - for example number of reads normalized by marker length and sequencing depth produces 'copies per million (CPMs), a quantitative estimate of abundance, and its flexible software architecture allows for its easy extension and custom reporting.

CORALE is capable of inferring presence of novel species

To demonstrate CORALE's ability to report sensible results when the provided reference does not contain a species producing a signal, we return to the hold-out set and remaining reference described before, and prepare an input of 338 samples each with a single 'novel' eukaryotic species at 0.1 genome coverage. We run CORALE, EukDetect, "EukDetect+": a version of EukDetect modified to filter MAPQ >= 5 instead of MAPQ >= 30, and "EukDetect-" with functionality simplified to report all taxa where least four reads align with MAPQ

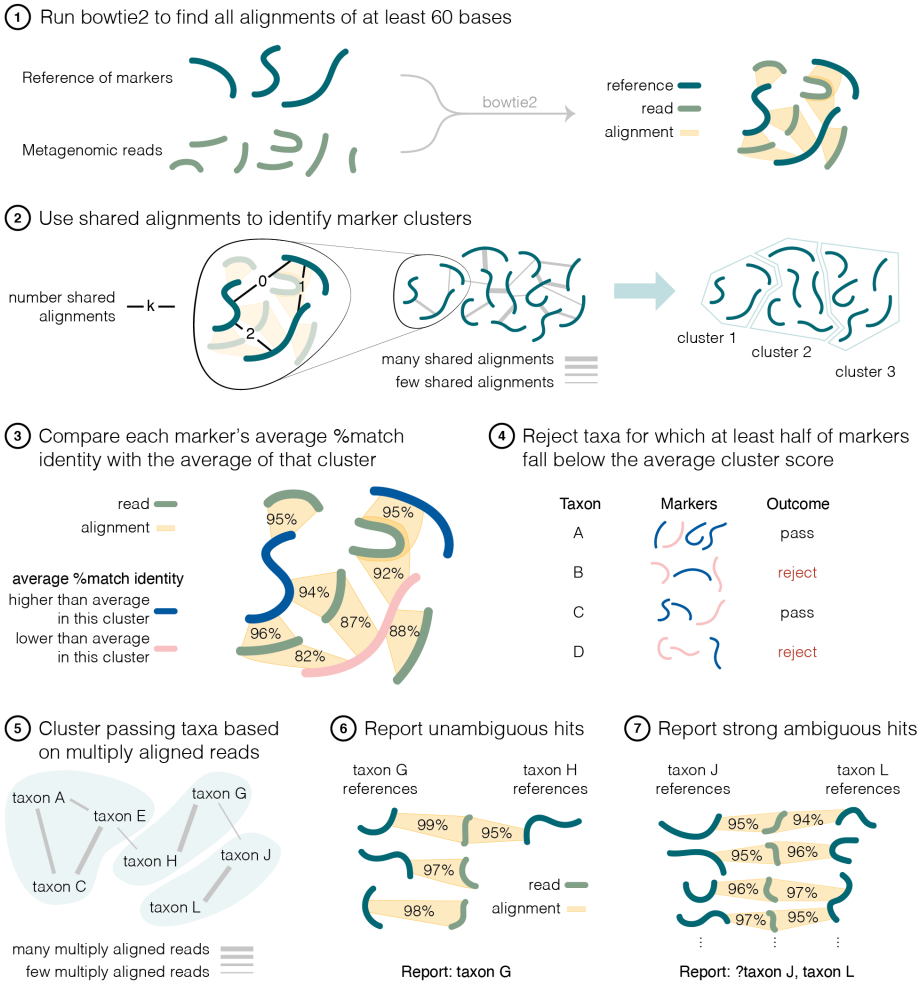

Figure 2: **CORALE - schematic**

>= 30 to at least two markers. We then classify results by correctness measured by taxonomic proximity - whether results reported were all of the same genus as the source species, as well as correctness measured by correct plurality of results.

CORALE does best out of the four tools at recognising presence of eukaryotic signal in the sample, as well as at reporting a single result in the correct genus (Figure 3). The MAPQ >= 30 filter makes inferring novel species more difficult, since EukDetect's proportion of No Hits is higher than for EukDetect+ and exactly the same as EukDetect-. Modifying EukDetect to filter on MAPQ >= 5 is not an adequate adjustment, because while it improves the tool's ability to recognise eukaryotic signal in the sample, it compromises the tool's ability to recognise that this signal consists of only a single species. Our method maintains the second ability while improving the first one - fractions of samples where signal is detected for CORALE, EukDetect, EukDetect+, and EukDetect - are respectively 0.607, 0.346, 0.47, and 0.346, and fractions of detected signal reported as one species are respectively 0.8, 0.812, 0.604, and 0.632.

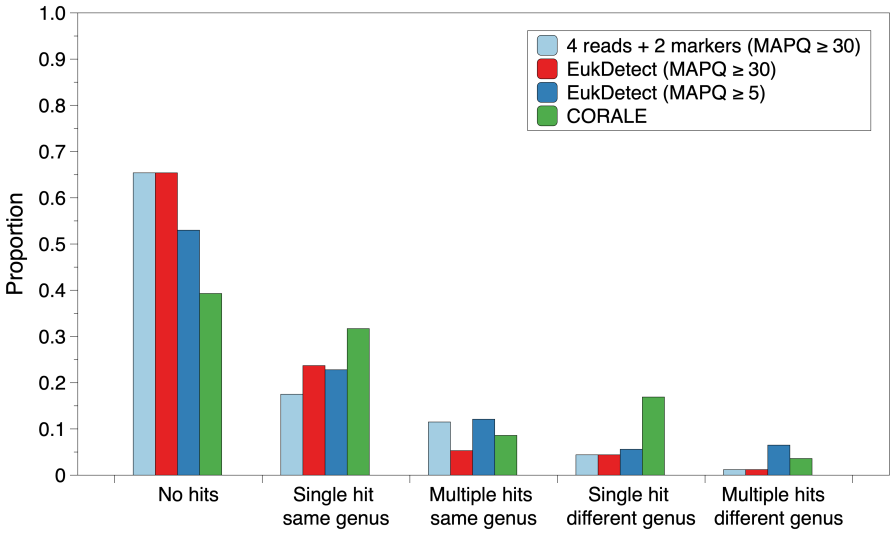

Figure 3: **Simulated unknown taxa - 0.1 coverage**

Evaluating CORALE on human microbiome data

To test our method on data where there are expectations about which eukaryotes might be present, we turn to the the DIABIMMUNE study (Vatanen et al. 2016), for which 136 data points about 30 different eukaryotes were reported across 1154 samples in the original EukDetect publication (Lind and Pollard 2021). Processing these same 1154 samples, CORALE is in exact concordance

with EukDetect on 122/136 data points, and adds additional 97 data points. CORALE reports common taxa at a higher frequency, for example, *S. cerevisiae* appears 67 times, while EukDetect only identifies this organism 31 times, suggesting higher sensitivity. No additional hit reported by CORALE is obviously implausible - they consist primarily of yeast and other fungi that have been previously reported in the human gut - and the lack of absurd results that happen when alignments to taxon-specific databases are used without caution (R Marcelino, Holmes, and Sorrell 2020) confirms specificity is not compromised.

Importantly, CORALE differs from EukDetect in how it treats reads that might originate from a novel species. For example, in sample G78909 from DIABIMMUNE, EukDetect reports *Penicillium nordicum*, while our method reports a novel *Penicillium*. In sample G80329, our method agrees with EukDetect regarding detection of *Candida parapsilosis*, and also identifies an additional *C. albicans*. Finally, in sample G78500 EukDetect reports *Saccharomyces cerevisiae* and *Kazachstania unispora*, which our method reports to be reads from one taxon slightly different from the reference *Saccharomyces cerevisiae*.

Automating eukaryote detection with CORALE

In addition to making our software freely available, we integrate CORALE into the automated data loading workflow for our open-science platform, MicrobiomeDB.org. Starting with Release 25 (2 Dec 2021), all whole genome sequencing data on the site is profiled for presence of eukaryotes using our approach (Figure 2). As of Release 25, the site contains 5113 samples from 6 published metagenomic studies Gasparrini et al. (2019), and CORALE identifies 97 distinct fungal species across 1661/5113 (32%) metagenomic samples. A summary of the top 10 most frequently observed fungi (Figure 3A) reveals that *Malassezia restricta*, a common commensal and opportunistic pathogen, and *Candida albicans*, a prevalent component of gut flora, are most commonly detected, with presence in 363 and 229 samples, respectively. Since these results are integrated with all other sample annotations on MicrobiomeDB, users can easily identify associations between eukaryotes and metadata (Figure 3B and 3C). For example, *Malassezia globosa* (identified in 130 samples) is primarily found on skin and nostrils, while *C. albicans* (Figure 3B and 3C), *C. parapsilosis*, *Clavispora lusitaniae*, and *S. cerevisiae* are all primarily or exclusively found in stool.

A

Top 10 fungal species detected

across all metagenomic studies on MicrobiomeDB

species	sample count
Malassezia restricta[1]	363
Candida albicans[2]	229
Purpureocillium lilacinum[3]	181
Saccharomyces cerevisiae[2]	145
Malassezia globosa[1]	130
Candida parapsilosis[2]	113
Clavispora lusitaniae[2]	60
Cyberlindnera jadinii	43
Candida tropicalis	32
Malassezia sympodialis[1]	17

[1] Primarily on skin and nostrils
[2] Primarily or exclusively in stool
[3] Exclusively in infant gut

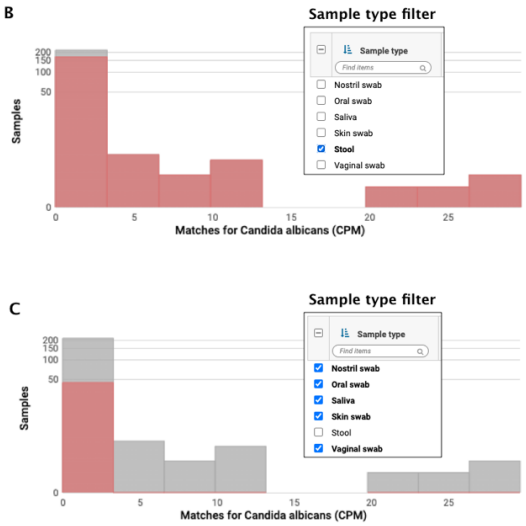

Figure 4: **MicrobiomeDB - top fungal species**

Discussion

Use cases and limitations of CORALE

CORALE is a tool for taxonomic profiling based on the same reference and with the same aim as EukDetect (Lind and Pollard 2021): to report on eukaryotes in metagenomes without overwhelming the results by false positives. Thanks to the novel approach involving the use of many alignments per read and Markov clustering, we achieve better sensitivity and alleviate reference bias associated with filtering on MAPQ scores. Adequate specificity combined with the tool's reliability lets us integrate eukaryotic detection into our open-science platform, MicrobiomeDB.org, and provides the microbiome research community with a means of setting up broad screens of metagenomic data to identify samples where eukaryotes are present.

Unfortunately, the approach of using shotgun metagenomics followed by processing with CORALE is not without limitations, chiefly due to high cost of WGS sequencing to the depth required to detect most eukaryotes. CORALE handles these limitations gracefully, making use of minimal information required to plausibly report unambiguous hits. Future improvements in genome assembly could be of help here through yielding better information on eukaryote-specific genomic sequences which could be used to create a larger reference with more taxa and more sequences per taxon, but CORALE does not require a complete reference - it is able to infer presence of taxa which are absent from the reference.

Future work

Our approach could potentially be applied to processing alignments to any reference that is anticipated to be redundant and incomplete, and where reads are expected to map with varying identity. This includes identification of bacteria to the strain-level resolution required in genomic epidemiology, as well as taxonomic classification of viral reads.

Additionally, the efficacy of using many alignments per read in combination with clustering demonstrated by CORALE shows that to develop new metagenomics tools, it can be worthwhile to view protein sequences as a similarity-based network: naturally occuring proteins form isolated clusters of similar sequence (Smith 1970). With additional theoretical work, this view could become a basis for predictions about presence of eukaryotes, potentially providing probabilistic estimates of certainty on reported results.

Conclusion

CORALE (for Clustering of Reference Alignments) is a tool for identification of eukaryotes in shotgun metagenomic studies. While CORALE is based on the same marker gene reference as EukDetect, it does not use EukDetect's approach to filtering, most notably not including the MAPQ >= 30 filter which we show to have species-specific impact on results. Its approach, based on multiple alignments and Markov clustering, results in sensitive and accurate detection, and is capable of inferring presence of eukaryotes not included in the reference. We show this to be the case using simulated samples with 'novel' species, as well as data from DIABIMMUNE, a large infant gut metagenome study. CORALE is also successfully deployed on our MicrobiomeDB.org resource, demonstrating the appropriateness of our method for large-scale screens of metagenomic data for the purpose of detecting eukaryotes.

Methods

Simulated reads

To conduct simulations, we use `wgsim` (Li 2011) to simulate reads from the 1/23/2021 version of EukDetect's reference, latest at time of writing, consisting of BUSCOs from OrthoDB (Kriventseva et al. 2019). We use `bowtie2` (Langmead and Salzberg 2012) to align reads to references, in end to end (default) mode. When using `wgsim` we set read length to 100, base error rate to 0, and other parameters set to their default values unless otherwise specified.

We process results using custom Python code. We judge correctness of simulated alignments using the level of the lowest common taxon containing source and match, as provided by the ETE toolkit (Huerta-Cepas, Serra, and Bork 2016) using the NCBI database version dated 2020/1/14 packaged with EukDetect. We deem the alignment correct if the source and match are of the same species, and in case of hold-out analysis when this is not possible by construction, same genus.

Our calculations of precision and recall are as defined in by OPAL framework (Meyer et al. 2019): precision is proportion of correctly mapping reads among reads that map to any reference, and recall is proportion of sampled reads that correctly map.

Simulated whole samples

To simulate whole samples, we skip 33 inputs where `wgsim` considers too fragmented to source reads from at a set coverage, yielding 338 samples. We select the number of reads to source per marker so as to achieve uniform 0.1 coverage, as calculated in (Sims et al. 2014).

To run EukDetect, we edit the default config file such that it lists the simulated samples. For the MAPQ >= 5 modification ("EukDetect+"), we modify source code of our local installation. For the "four reads mapping with at MAPQ >= 30 to at least two markers" modification, ("EukDetect-"), we run CORALE in non-default configuration use these three filters.

Data availability

All our software is publicly available under the MIT license: CORALE (*github.com/wbazant/CORALE*), its main Python module, (*github.com/wbazant/marker_alignments*), and a mix of Python, Make, and Bash scripts to produce simulations, comparisons, and figures for this publication (*github.com/wbazant/markerAlignmentsPaper*).

All results are publicly viewable and downloadable on MicrobiomeDB. In addition, the following files are available as supplemental material:

Simulated whole samples - results for different methods

Simulated reads - per-species breakdown and aggregate stats

DIABIMMUNE - CORALE vs EukDetect comparison

#Bibliography

"Aspergillus Fumigatus and Aspergillosis." 1999. *Clinical Microbiology Reviews* 12 (2): 310–50. https://doi.org/10.1128/cmr.12.2.310.

Beghini, Francesco, Edoardo Pasolli, Tin Duy Truong, Lorenza Putignani, Simone M Cacciò, and Nicola Segata. 2017. "Large-Scale Comparative Metagenomics of Blastocystis, a Common Member of the Human Gut Microbiome." *The ISME Journal* 11 (12): 2848–63.

Breitwieser, Florian P, DN Baker, and Steven L Salzberg. 2018. "KrakenUniq: Confident and Fast Metagenomics Classification Using Unique k-Mer Counts." *Genome Biology* 19 (1): 1–10.

Burrows, Michael, and David Wheeler. 1994. "A Block-Sorting Lossless Data Compression Algorithm." In *Digital SRC Research Report*. Citeseer.

Consortium, Human Microbiome Project et al. 2012. "Structure, Function and Diversity of the Healthy Human Microbiome." *Nature* 486 (7402): 207.

Doron, Itai, Irina Leonardi, Xin V. Li, William D. Fiers, Alexa Semon, Meghan Bialt-DeCelie, Mélanie Migaud, et al. 2021. "Human Gut Mycobiota Tune Immunity via CARD9-Dependent Induction of Anti-Fungal IgG Antibodies." *Cell* 184 (4): 1017–1031.e14. https://doi.org/10.1016/j.cell.2021.01.016.

Doron, Itai, Marissa Mesko, Xin V Li, Takato Kusakabe, Irina Leonardi, Dustin G Shaw, William D Fiers, et al. 2021. "Mycobiota-Induced IgA Antibodies Regulate Fungal Commensalism in the Gut and Are Dysregulated in Crohn's Disease." *Nature Microbiology* 6 (12): 1493–1504.

Gasparrini, Andrew J., Bin Wang, Xiaoqing Sun, Elizabeth A. Kennedy, Ariel Hernandez-Leyva, I. Malick Ndao, Phillip I. Tarr, Barbara B. Warner, and Gautam Dantas. 2019. "Persistent Metagenomic Signatures of Early-Life Hospitalization and Antibiotic Treatment in the Infant Gut Microbiota and Resistome." *Nature Microbiology* 4 (12): 2285–97. https://doi.org/10.1038/s41564-019-0550-2.

Gibson, Molly K., Bin Wang, Sara Ahmadi, Carey-Ann D. Burnham, Phillip I. Tarr, Barbara B. Warner, and Gautam Dantas. 2016. "Developmental Dynamics of the Preterm Infant Gut Microbiota and Antibiotic Resistome." *Nature Microbiology* 1 (4). https://doi.org/10.1038/nmicrobiol.2016.24.

Hayden, Hillary S., Alexander Eng, Christopher E. Pope, Mitchell J. Brittnacher, Anh T. Vo, Eli J. Weiss, Kyle R. Hager, et al. 2020. "Fecal Dysbiosis in Infants with Cystic Fibrosis Is Associated with Early Linear Growth Failure." *Nature Medicine* 26 (2): 215–21. https://doi.org/10.1038/s41591-019-0714-x.

Huerta-Cepas, Jaime, François Serra, and Peer Bork. 2016. "ETE 3: Reconstruction, Analysis, and Visualization of Phylogenomic Data." *Molecular Biology and Evolution* 33 (6): 1635–38.

Jiang, Tony T., Tzu-Yu Shao, W. X. Gladys Ang, Jeremy M. Kinder, Lucien H. Turner, Giang Pham, Jordan Whitt, Theresa Alenghat, and Sing Sing Way. 2017. "Commensal Fungi Recapitulate the Protective Benefits of Intestinal Bacteria." *Cell Host & Microbe* 22 (6): 809–816.e4. https://doi.org/10.1016/j.chom.2017.10.013.

Kostic, Aleksandar D., Dirk Gevers, Heli Siljander, Tommi Vatanen, Tuulia

Hyötyläinen, Anu-Maaria Hämäläinen, Aleksandr Peet, et al. 2015. "The Dynamics of the Human Infant Gut Microbiome in Development and in Progression Toward Type 1 Diabetes." *Cell Host & Microbe* 17 (2): 260–73. https://doi.org/10.1016/j.chom.2015.01.001.

Kriventseva, Evgenia V, Dmitry Kuznetsov, Fredrik Tegenfeldt, Mosè Manni, Renata Dias, Felipe A Simão, and Evgeny M Zdobnov. 2019. "OrthoDB V10: Sampling the Diversity of Animal, Plant, Fungal, Protist, Bacterial and Viral Genomes for Evolutionary and Functional Annotations of Orthologs." *Nucleic Acids Research* 47 (D1): D807–11.

Laforest-Lapointe, Isabelle, and Marie-Claire Arrieta. 2018. "Microbial Eukaryotes: A Missing Link in Gut Microbiome Studies." *MSystems* 3 (2): e00201–17.

Langmead, Ben, and Steven L Salzberg. 2012. "Fast Gapped-Read Alignment with Bowtie 2." *Nature Methods* 9 (4): 357–59.

Leonardi, Irina, Iris H. Gao, Woan-Yu Lin, Megan Allen, Xin V. Li, William D. Fiers, Meghan Bialt De Celie, et al. 2022. "Mucosal Fungi Promote Gut Barrier Function and Social Behavior via Type 17 Immunity." *Cell*. https://doi.org/10.1016/j.cell.2022.01.017.

Li, Heng. 2011. "Wgsim-Read Simulator for Next Generation Sequencing." *Github Repository*.

Lind, Abigail L, and Katherine S Pollard. 2021. "Accurate and Sensitive Detection of Microbial Eukaryotes from Whole Metagenome Shotgun Sequencing." *Microbiome* 9 (1): 1–18.

Meyer, Fernando, Andreas Bremges, Peter Belmann, Stefan Janssen, Alice C McHardy, and David Koslicki. 2019. "Assessing Taxonomic Metagenome Profilers with OPAL." *Genome Biology* 20 (1): 1–10.

Mora, Camilo, Derek P. Tittensor, Sina Adl, Alastair G. B. Simpson, and Boris Worm. 2011. "How Many Species Are There on Earth and in the Ocean?" Edited by Georgina M. Mace. *PLoS Biology* 9 (8): e1001127. https://doi.org/10.1371/journal.pbio.1001127.

Nash, Andrea K, Thomas A Auchtung, Matthew C Wong, Daniel P Smith, Jonathan R Gesell, Matthew C Ross, Christopher J Stewart, et al. 2017. "The Gut Mycobiome of the Human Microbiome Project Healthy Cohort." *Microbiome* 5 (1): 1–13.

Oliveira, Francislon S, John Brestelli, Shon Cade, Jie Zheng, John Iodice, Steve Fischer, Cristina Aurrecoechea, et al. 2018. "MicrobiomeDB: A Systems Biology Platform for Integrating, Mining and Analyzing Microbiome Experiments." *Nucleic Acids Research* 46 (D1): D684–91.

Olm, Matthew R., Nicholas Bhattacharya, Alexander Crits-Christoph, Brian A. Firek, Robyn Baker, Yun S. Song, Michael J. Morowitz, and Jillian F. Banfield. 2019. "Necrotizing Enterocolitis Is Preceded by Increased Gut Bacterial Replication, Klebsiella , and Fimbriae-Encoding Bacteria." *Science Advances* 5 (12). https://doi.org/10.1126/sciadv.aax5727.

Ost, Kyla S., Teresa R. O'Meara, W. Zac Stephens, Tyson Chiaro, Haoyang Zhou, Jourdan Penman, Rickesha Bell, et al. 2021. "Adaptive Immunity Induces Mutualism Between Commensal Eukaryotes." *Nature* 596 (7870):

114–18. https://doi.org/10.1038/s41586-021-03722-w.

R Marcelino, Vanessa, Edward C Holmes, and Tania C Sorrell. 2020. "The Use of Taxon-Specific Reference Databases Compromises Metagenomic Classification." *BMC Genomics* 21 (1): 1–5.

Schoch, Conrad L. et al. 2012. "Nuclear Ribosomal Internal Transcribed Spacer (ITS) Region as a Universal DNA Barcode Marker for fungi." *Proceedings of the National Academy of Sciences* 109 (16): 6241–46. https://doi.org/10.1073/pnas.1117018109.

Sims, David, Ian Sudbery, Nicholas E Ilott, Andreas Heger, and Chris P Ponting. 2014. "Sequencing Depth and Coverage: Key Considerations in Genomic Analyses." *Nature Reviews Genetics* 15 (2): 121–32.

Smith, John Maynard. 1970. "Natural Selection and the Concept of a Protein Space." *Nature* 225 (5232): 563–64.

Thankaswamy-Kosalai, Subazini, Partho Sen, and Intawat Nookaew. 2017. "Evaluation and Assessment of Read-Mapping by Multiple Next-Generation Sequencing Aligners Based on Genome-Wide Characteristics." *Genomics* 109 (3-4): 186–91.

Vatanen, Tommi, Aleksandar D Kostic, Eva d'Hennezel, Heli Siljander, Eric A Franzosa, Moran Yassour, Raivo Kolde, et al. 2016. "Variation in Microbiome LPS Immunogenicity Contributes to Autoimmunity in Humans." *Cell* 165 (4): 842–53.

Wibbelt, Gudrun, Andreas Kurth, David Hellmann, Manfred Weishaar, Alex Barlow, Michael Veith, Julia Prüger, et al. 2010. "White-Nose Syndrome Fungus (Geomyces Destructans) in Bats, Europe." *Emerging Infectious Diseases* 16 (8): 1237–43. https://doi.org/10.3201/eid1608.100002.

Wilson, Richard A., and Nicholas J. Talbot. 2009. "Under Pressure: Investigating the Biology of Plant Infection by Magnaporthe Oryzae." *Nature Reviews Microbiology* 7 (3): 185–95. https://doi.org/10.1038/nrmicro2032.